

Limits of Deep Learning

Simon Jenni
University of Bern
Email: simujenni@students.unibe.ch

Simon Brulhart
University of Fribourg
Email: simon.brulhart@unifr.ch

Abstract—The abstract goes here.

I. INTRODUCTION

Recent state of the art natural language models represent words by embedding them as vectors in a continuous vector space [1] [2]. These embeddings are learned as weights of a recurrent neural network language model. It has been demonstrated that the distributed representation of words as vectors in a common vector-space captures not just syntactic similarities (e.g. *cat* is close to *cats*) but semantic similarities as well. Concretely, the model's ability to perform word-arithmetic has been demonstrated [3]. To exemplify this, let f be the word-embedding function. Ideally, we would then have an embedding where distances in vector-space give a measure of semantic-similarity between the words. For two word pairs such as (*king* : *man*) and (*queen* : *woman*) which have the same relation (or similarity) we would obtain $f(\text{king}) - f(\text{man}) = f(\text{queen}) - f(\text{woman})$ or equivalently $f(\text{king}) - f(\text{man}) + f(\text{woman}) = f(\text{queen})$. This of course is equivalent to forming the analogy "*king* is to *man* as *queen* is to *woman*".

The goal of this work is to further evaluate the performance of the popular word2vec tool on these tasks. We evaluate the effect of different model and learning parameters on the performance and test the limits of the models semantic-analysis capabilities on new challenging test data.

A. Related Work

The work by Tomas Mikolov [1], where the models were introduced, reports results on a test set they themselves created. The set contained five types of semantic questions where relations such as *Country* ~ *Currency* or *Country* ~ *Capitol* were tested. In [3] the models performance on semantical similarity has been demonstrated on SemEval-2012 Task 2 [4], where the objective is to decide how similar two target word-pairs such as (*glass* : *break*) and (*soldier* : *fight*) are with respect to a relation "*anXtypicallyY*".

B. Models of word2vec

word2vec is an implementation of the Continuous Bag of Words (CBOW) and Skip-Gram models introduced in [1]. Both of these models are based on shallow two-layer neural networks with architectures depicted in Figure 1 and Figure 2. Both architectures take one-hot encoded words as inputs which get then mapped into an N -dimensional vector via the weight matrix W . The hidden unit then accumulates the sum

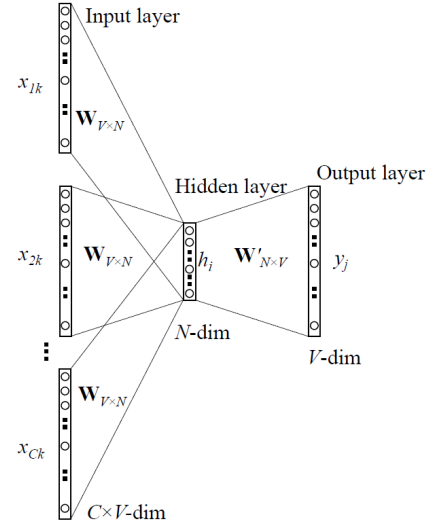


Fig. 1: Architecture of the CBOW model.

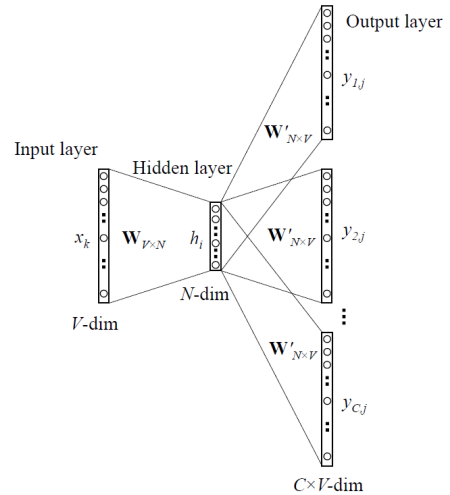


Fig. 2: Architecture of the Skip-Gram model.

of these vectors for all input words. The output of the model is then generated multiplying the hidden-units by weights W' and applying soft-max to the result in order to generate a valid probability-distribution over words in the vocabulary.

The models are learnt in an unsupervised fashion where either the centre-word of a series of words is predicted from

its context (CBOW) or the context is predicted from the centre-word (Skip-Gram). Both models try to maximize the likelihood of their output. Training is done via gradient descent using back-propagation [5] to compute derivatives.

II. EXPERIMENTS

A. Setup

B. New Test Data

III. CONCLUSION

The conclusion goes here.

ACKNOWLEDGMENT

The authors would like to thank...

REFERENCES

- [1] T. Mikolov, K. Chen, G. Corrado, and J. Dean, "Efficient estimation of word representations in vector space," *arXiv preprint arXiv:1301.3781*, 2013.
- [2] J. Pennington, R. Socher, and C. D. Manning, "Glove: Global vectors for word representation," in *EMNLP*, vol. 14, 2014, pp. 1532–1543.
- [3] T. Mikolov, W.-t. Yih, and G. Zweig, "Linguistic regularities in continuous space word representations," in *HLT-NAACL*, 2013, pp. 746–751.
- [4] D. A. Jurgens, P. D. Turney, S. M. Mohammad, and K. J. Holyoak, "Semeval-2012 task 2: Measuring degrees of relational similarity," in *Proceedings of the First Joint Conference on Lexical and Computational Semantics-Volume 1: Proceedings of the main conference and the shared task, and Volume 2: Proceedings of the Sixth International Workshop on Semantic Evaluation*. Association for Computational Linguistics, 2012, pp. 356–364.
- [5] D. E. Rumelhart, G. E. Hinton, and R. J. Williams, "Learning representations by back-propagating errors," *Cognitive modeling*, vol. 5, no. 3, p. 1, 1988.